

Bruno Marcuche

Platform Architect

Boulder, CO 80301 · bmarcuche@gmail.com · linkedin.com/in/bruno-marcuche · resume.mindtunnel.org · github.com/bmarcuche



Summary

Platform Architect and technical leader who builds the systems other teams run on. I've scaled infrastructure across on-prem, hybrid, and cloud, and automated deployment and operations for thousands of Linux and Windows instances. Most recently I architected an internal AI agent platform: a local semantic router routes 85% of ops requests without LLM reasoning, and median upgrade lead time fell from 40 to 11 days. I lead ops and SRE teams, drive observability with OpenTelemetry and PagerDuty, and turn slow, manual operations into fast, repeatable automation.



Professional Experience

08/2026 to present
Boulder, CO

Platform Architect AssetWorks

- Architect and operate the hosting platform behind a large fleet of managed customer environments, applying AI-driven engineering to fleet operations at scale.
- Built Sentry, agentless monitoring that checks every hosted database every 2 to 10 minutes, so problems surface before they cause an outage: in its first 60 days it flagged 65 databases with archive-log errors and 62 crashed services that would not recover on their own.
- Led GhostWatch, an AI-assisted incident pipeline from the Microsoft Global Hackathon 2026. Its detector reads report-server behavior, not just uptime, so in replay it caught 92% of hangs a median 107 minutes early at a 4% false alarm rate; grouping alerts into incidents cut on-call triage 2.7x.
- Added self-healing on top of the environment inventory, so 354 outages were repaired automatically.
- Lead capacity and sizing analysis for executive decisions, and root-cause investigation on production incidents spanning application servers, storage, and reporting services.

03/2025 to 07/2026
Berwyn, PA

Operations Team Lead AssetWorks

- Built an internal AI agent platform on the Model Context Protocol with a local semantic router in front of every ops request, so 85% route without the LLM reasoning fallback at zero LLM tokens per routing decision.

- Gave agents typed tools and a shared memory, so 98.6% of routed tasks complete; guardrail hooks stopped 211 destructive deletes and 117 credential exposures before they ran.
- Moved customer upgrades onto agent-built pipelines, so median upgrade lead time fell from 40 to 11 days and operator wait between stages dropped 89%.
- Built the Hosted Environment Navigator: continuous discovery replaced remoting into servers one at a time, so any environment is one search away and 99% of the inventory maintains itself.
- Built Hosted Access Manager: self-service, time-boxed grants replaced DBA tickets, so access takes about a minute instead of a day, privileged accounts stay locked 99.6% of the time, and every grant closes with zero DBA cleanup.
- Led a team of five; rolled out OpenTelemetry and Observe; mentored through 1:1s and training.

12/2022 to 01/2025
Boulder, CO

Backend Developer, Founder
EdventureTrek

- Founded an educational exploration game; designed custom taxonomy GPTs for plant and animal classification.
- Built the Python/FastAPI backend with MySQL and event logging; ran CI/CD on GCP with GitHub Actions.

07/2022 to 12/2022
Atlanta, GA

Site Reliability Engineering Manager
AnswerRocket

- Led a remote SRE team of four; ran weekly syncs and architecture reviews.
- Expanded Ansible coverage across AWS (15% less manual deploy time) and automated SOC 2 environment validation.

03/2016 to 07/2022
Alpharetta, GA

Site Reliability Architect
OfficeSpace Software

- Owned production and staging on GCP: patching, config management, release packaging, and automation with Puppet and Python.
- Migrated from Rackspace to GCP, saving \$60K annually; cut deploy times over 60% with a CircleCI, Puppet, Docker and Terraform pipeline.
- Built a Slackbot that deploys customer instances in under 10 minutes; hired, onboarded and led a three-person SRE team.

2009 to 2016
São Paulo, Brazil

Sr. Technical Consultant / Team Lead
Hewlett Packard

- Tier 3 support for HP Server Automation; automated workflows with Python against the HPSA API.
- Mentored junior engineers; ranked #1 in the team for customer satisfaction.



Education

10/2001 to 12/2004
Ft. Lauderdale, Florida

Information Systems | Bachelor's Degree
ITT Technical Institute
(Honors Graduate)



Strengths

Leadership AI-Led Ops Cloud Infrastructure Automation



Key Skills

CLOUD & AUTOMATION

GCP Azure Cloud Cloud Run Terraform Ansible Docker

CI/CD & DELIVERY

GitHub Actions Azure DevOps Jenkins

OBSERVABILITY

OpenTelemetry Prometheus PagerDuty Observe

AI & AGENTS

LLM Agents Model Context Protocol pgvector Hugging Face Claude

DEVELOPMENT & DATA

Python Redis Postgres

SYSTEMS & SERVING

Linux Nginx



Hobbies

Exploring AI tooling (LLMs, MCP) SRE meetups Home lab experimentation with Docker



Volunteering

07/2020 to 11/2024
Boulder, Colorado, USA

Delivery Driver / Wellness Check
Meals on Wheels Boulder

As a Meals on Wheels delivery driver, I got to enjoy great conversations with some of Boulder's greatest citizens.
mowboulder.org