

Bruno Marcuche

Platform Architect

Boulder, CO 80301 · bmarcuche@gmail.com · linkedin.com/in/bruno-marcuche · resume.mindtunnel.org · github.com/bmarcuche

Summary

Platform Architect and technical leader who builds the systems other teams run on. I've scaled infrastructure across on-prem, hybrid, and cloud, and automated deployment and operations for thousands of Linux and Windows instances. Most recently I architected an internal AI agent platform, a custom semantic router orchestrating specialized LLM agents, that has handled 10,000+ ops and engineering tasks and cut change lead time by ~89%. I lead ops and SRE teams, drive observability with OpenTelemetry and PagerDuty, and turn slow, manual operations into fast, repeatable automation.

Professional Experience

08/2026 to present
Boulder, CO

Platform Architect AssetWorks

- Architect and operate the hosting platform behind a large fleet of managed customer environments, applying AI-driven engineering to fleet operations at scale.
- Own platform architecture spanning Windows server fleets, Azure DevOps pipelines, and multi-tenant upgrade orchestration, plus relational schema design and maintenance.
- Built and run production agentic automation: a Python orchestration layer driving Claude with bi-encoder/cross-encoder retrieval over pgvector, feeding live operational systems.
- Design and build the internal tooling the hosting org runs on: fleet-wide database monitoring, upgrade scheduling with ticketing and calendar integration, and cross-system customer reconciliation.
- Maintain the environment inventory as the authoritative source of truth, with automated discovery and configuration-drift detection across the fleet.
- Lead capacity and sizing analysis for executive decisions, and root-cause investigation on production incidents spanning application servers, storage, and reporting services.

03/2025 to 07/2026
Berwyn, PA

Operations Team Lead AssetWorks

- Architected and built an internal AI agent platform on the Model Context Protocol: specialized LLM agents coordinated by a custom semantic router (fine-tuned sentence-transformer embeddings with pgvector knowledge retrieval). It handled 10,000+ routed ops and engineering tasks in its first 12 weeks across deployment, cloud, CI/CD, and incident response.
- Cut change lead time by ~89% after launch. Customer upgrades dropped from ~27 days to under 3 days, and provisioning of new environments from ~12 days to ~1.5 days, declining every month after going live (DORA lead time for changes).
- Delivered 235 Ansible and CI/CD pipelines built by the agents and automated 213 customer upgrade deployments, eliminating ~426 hours of manual deploy work, at 99.99% platform uptime.
- Designed multi-agent incident workflows that investigate and remediate across the fleet in a single session. They caught a bug that was deleting configuration fleet-wide and identified the root cause of an OS update and API regression that would otherwise take hours of manual log correlation.
- Lead a team of five; drove an observability rollout (OpenTelemetry and Observe) to reduce MTTR and mentor engineers through 1:1s, training, and knowledge sharing.

12/2022 to 01/2025
Boulder, CO

Backend Developer, Founder
EdventureTrek

- Founded and led development of an educational game focused on outdoor exploration and biodiversity.
- Designed custom taxonomy GPTs for plant and animal classification.
- Built Python/FastAPI backend with MySQL and event logging.
- Managed CI/CD on GCP with GitHub Actions and internal tooling.

07/2022 to 12/2022
Atlanta, GA

Site Reliability Engineering Manager
AnswerRocket

- Led remote SRE team (4 reports); ran weekly syncs and architecture reviews.
- Expanded Ansible coverage across AWS, cutting manual deploy time by 15%.
- Supported SOC 2 audit by automating cloud environment validation.

03/2016 to 07/2022
Alpharetta, GA

Site Reliability Architect
OfficeSpace Software

- Led SRE hiring, onboarding and 1:1s for a 3-person team.
- Built Slackbot enabling teams to deploy customer instances in under 10 minutes.
- Reduced deploy times over 60% via CI pipeline (CircleCI, Puppet, Docker, Terraform).

- Migrated infrastructure from Rackspace to GCP, saving \$60K annually.
- Owned production and staging infrastructure on GCP; managed OS patching, config management, release packaging, and automation with Puppet and Python.

2009 to 2016
São Paulo, Brazil

Sr. Technical Consultant / Team Lead
Hewlett Packard

- Delivered Tier 3 support for HP Server Automation; mentored junior engineers.
- Automated workflows by developing Python scripts utilizing the HPSA API.
- Ranked #1 in team for customer satisfaction.

 Education

10/2001 to 12/2004
Ft. Lauderdale, Florida

Information Systems | Bachelor's Degree
ITT Technical Institute
(Honors Graduate)

 Strengths

Leadership	AI-Led Ops
Cloud Infrastructure	Automation

 Key Skills

CLOUD & AUTOMATION

GCP	Azure Cloud
Cloud Run	Terraform
Ansible	Docker

CI/CD & DELIVERY

GitHub Actions	Azure DevOps
Jenkins	

OBSERVABILITY

OpenTelemetry	Prometheus
PagerDuty	Observe

AI & AGENTS

LLM Agents	Model Context Protocol
pgvector	Hugging Face
Claude	

DEVELOPMENT & DATA

Python
Postgres

Redis

SYSTEMS & SERVING

Linux

Nginx

✧ Hobbies

Exploring AI tooling (LLMs, MCP)

SRE meetups

Home lab experimentation with
Docker

♥ Volunteering

07/2020 to 11/2024
Boulder, Colorado, USA

Delivery Driver / Wellness Check Meals on Wheels Boulder

As a Meals on Wheels delivery driver, I got to enjoy great conversations with some of Boulder's greatest citizens.

mowboulder.org